Scalable Logic Synthesis using a Simple Circuit Structure

Alan Mishchenko   Robert Brayton

EECS Department, University of California, Berkeley, CA 94720
{alanmi, brayton}@eecs.berkeley.edu

Abstract

This paper proposes an alternate approach to logic synthesis using rewriting and peephole optimization but from a modern perspective. We use a simple logic structure (AIGs) as the basis for all the algorithms, and rely on efficient techniques, such as precomputation, reconvergence analysis, cut enumeration, Boolean matching, exhaustive simulation of small logic cones, and local resource-aware decision procedures based on Boolean satisfiability. The result is a logic synthesis flow that is orders of magnitude faster than traditional ones and more scalable, being applicable to large industrial netlists with millions of gates.

1 Introduction

Historically, logic synthesis approaches can be divided into two categories, local rule-based transformations (or rewriting) and technology independent/dependent algorithms.

Rewriting. This was used by some of the first logic synthesis methods [11]. In these, optimization was applied to a netlist composed of the gate types to be used in the final implementation. Then, so-called *peephole optimization* was applied, where the structure of a set of gates in a small window was recognized and associated with a set of transformations that could improve area or delay. The decision to "fire" a local transformation and the order of firing was controlled by a set of rules. The combinations of such methods became known as *rule-based logic synthesis*.

Some problems with this approach were:

- when new gates were added to the database, new rules were needed to control where and when such gates were to be used,
- as the rule-base grew, new rules interacted and interfered with the old rules, and managing the rule base became problematic,
- the rule base was specific to the set of gates being used, and
- the system became slow as the rule-base grew.

Algorithmic. Rule-based rewriting gave way to technology independent optimization [6][7] followed by technology dependent mapping [14]. The underlying notion was that there were a set of operations relatively good for a netlist, regardless of the final set of gates that would be used for the implementation. This allowed the synthesis problem to be decomposed into separate steps, both of which became amenable to deeper mathematical analysis, leading to an algorithmic treatment of sub-problems and improved understanding.

This required an abstract measure of goodness to guide the technology independent transformations. The *number of literals in the factored forms* became the standard metric for area, and many optimization algorithms were developed which were guided by improving this metric. For example, the operation "eliminate" collapses a node into its fanouts if the worth of a node computed using this metric did not exceed a specified threshold. "Kerneling"

selected divisors similarly. Node optimization minimized the logic function of a node represented in SOP form, which then was factored to obtain a literal count. Even though a minimum SOP and a factor are not unique, the heuristic used was that a good factorization of a minimum SOP was close enough to the best factorization possible for the node function. Thus, if the metric improved, the result was accepted and the current node function was replaced. Similar metrics were developed for delay. At the same time, algorithmic technology-mapping methods were developed and were steadily improved, becoming more sophisticated and complex.

Scalability required that logic synthesis algorithms be extended, or replaced by more scalable ones. However, this resulted in several disadvantages. The new algorithms

- increased the barrier to understanding,
- made implementation more difficult,
- required more complex code structures,
- made code maintenance more difficult, and
- were rarely made public, since much of these developments were done within commercial CAD tools becoming valuable assets of the industrial companies.

Alternative Approach. The contribution of this paper is an alternate approach to logic synthesis that is simple, fast, and scalable. We keep the separation between the technology independent and dependent aspects. Local AIG-based transformations, such as *rewriting* [2][20], *resubstitution*, and *redundancy removal*, are developed in a simplified form. Speed and scalability are achieved as follows:

- The network is always represented using a simple structure, an And-Inverter Graph (AIG), and all transformations work on this uniform representation. The metrics used for area and delay, the total number of AND nodes and depth of the graph, are easy to compute and have a direct impact on technology mapping (which uses the AIG as the subject graph).
- Instead of graph isomorphic and SOP-based optimization, Boolean functions of an AIG node are computed as functions of its various fanin cut variables. Cut computations use a recent development that can efficiently enumerate all cuts up to 12 inputs [9]. A quality/runtime tradeoff is controlled by selecting cuts of appropriate size.
- Boolean functions of cuts are represented by truth tables, which are used to hash into a database where new rewriting structures are found. Due to the high speed of bit-parallel manipulation, sub-problems arising in local optimization can be solved quickly by exhaustive simulation. This scales well for up to about 16 inputs and often works better than BDDs or SAT.

Since these types of local optimizations are fast, they can be repeated many times. For example, performing 10 rewriting passes over a typical network is still an order of magnitude faster than running traditional (but resource aware) synthesis in MVSIS, and several orders of magnitude faster than *script.alegraic* in SIS. In addition, the cumulative effect of these multiple optimization passes is often superior in quality.

The paper is organized as follows. Section 2 surveys Boolean networks, AIGs, windowing, and reconvergence-driven cut computation. Section 3 presents the algorithms, for local rewriting, resubstitution, and redundancy removal, which have been adapted for AIGs and made efficient by rigorous enforcement of resource limits. Section 4 reports experimental results of the new logic synthesis flow compared with traditional approaches. Section 5 concludes and lists directions for future work.

2 Background

2.1 Boolean networks

Definition. A *Boolean function* is a mapping from n-dimensional ($n \geq 0$) Boolean space into a 1-dimensional one: $\{0,1\}^n \to \{0,1\}$.

Definition. A *Boolean network* is a directed acyclic graph (DAG) with nodes represented by Boolean functions. The sources of the graph are the *primary inputs* (PIs) of the network; the sinks are the *primary outputs* (POs).

Definition. The output of a node may be an input to other nodes called its *fanouts*. The inputs of a node are called its *fanins*. If there is a path from node a to b, then a is in the *transitive fanin* of b and b in the *transitive fanout* of a. The transitive fanin of b, $TFI(b)$, includes node b and the nodes in its transitive fanin, including the PIs. The *transitive fanout* of b, $TFO(b)$, includes node b and the nodes in its transitive fanout, including the POs. An *edge* connects two nodes that are in the fanin/fanout relationship. The fanin/fanout of an edge is the fanin/fanout node of the connected pair of nodes.

Definition. A *maximum fanout free cone* (MFFC) of node n is a subset of $TFI(n)$, such that every path from a node in the subset to the POs passes through n. Informally, the MFFC of a node contains the node and all the logic used exclusively by the node. When a node is removed (substituted), the logic in its MFFC can also be removed.

Definition. A *cut C* of node n is a set of nodes of the network, called *leaves*, such that each path from a PI to n passes through at least one leaf. Node n is called *root* of cut C. The cut *size* is the number of its leaves. A *trivial cut* of the node is the cut composed of the node itself. A cut is *K-feasible* if the number of nodes in the cut does not exceed K. A cut is said to be *dominated* if there is another cut of the same node, which is contained, set-theoretically, in the given cut. The *volume* of a cut is the total number of nodes encountered on all paths between node n and the cut leaves.

Definition. The *local* function of an AIG node n, denoted $f_n(x)$, is a Boolean function of the logic cone rooted in n and expressed in terms of the leaves, x, which form a cut of n. The *global function* of an AIG node is its function in terms of the PIs of the network.

Definition. *Exhaustive simulation* is a practical way of checking equivalence of Boolean function of a node. Exhaustive simulation is performed using bitwise simulation of the cone with 2^k different input patterns, where k is the number of leaves (in practice k does not exceed 16). Another way of looking at exhaustive simulation is that it computes the truth-table of the root node in terms of the elementary truth-tables set at the leaves of some cut of this node.

2.2 And-Interver Graphs

And-Interver Graph (AIG) is a Boolean network composed of two-input ANDs and inverters. AIGs were used in a variety of applications since the early 60's [12] as a convenient representation for combinational logic of an arbitrary Boolean network. To derive an AIG, the SOPs of the nodes in the network are factored [6], the AND gates and OR gates of the factored forms are converted into two-input ANDs and inverters using DeMorgan's rule, and these nodes are added to the AIG manager.

Structural hashing of netlists composed of arbitrary gates was proposed and used in early CAD tools [11] to detect and merge isomorphic circuit structures. For AIGs, structural hashing is performed by one hash-table lookup when AND nodes are added to the AIG manager. It ensures that, for each pair of nodes, there is only one AND node having them as fanins (up to permutation).

Additionally, the AIG derived from a logic network is often *balanced* to reduce the number of AIG levels, by applying the associative transform, $a(bc) = (ab)c$. Both structural hashing and balancing are performed in one topological sweep from the PIs and have linear complexity in the number of AIG nodes.

The *size* (*area*) of the AIG is the number of its nodes. The *depth* (*delay*) of the AIG is the number of nodes on the longest path from the PIs to the POs. The goal of optimization by local transformations is to reduce both area and delay of an AIG, although in this paper we focus on the area optimization. Delay-oriented synthesis additionally makes the transformations aware of the AIG depth.

Software implementation of an AIG package is similar to that of an efficient BDD package [4]. Inverters are represented as flipped pointers to the AIG nodes. The AIG nodes have reference counters, which show the number of fanouts of each node. Reference counting leads to fast counting of nodes in an MFFC and efficient add/remove operations for individual nodes and their MFFCs.

Both BDD and AIG packages support the unique table to ensure that there is only one node with the given fanins. Due to the uniqueness of the Shannon expansion with respect to a variable, this leads to the canonical BDD structure for a given variable order. For AIGs, the AND-decomposition is not unique, so the use of a unique table guarantees only structural canonicity within one logic level (described above as structural hashing).

Maintaining an analogy with Reduced Ordered BDDs led to the development of the FRAIG package [18], which uses a balanced combination of simulation and SAT [16] to enforce functional canonicity of AIG nodes on-the-fly, as they are added to the package. The FRAIG package has found extensive use in logic synthesis, technology mapping [8], and equivalence checking [22] in ABC [1].

It should be noted that, although AIGs are used as the main representation in this paper and in ABC, other types of simple circuit structures would also work: NAND graphs, OR-INV graphs, AND-XOR-INV graphs, Reduced Boolean Circuits [3], etc. All these representations have similar size for practical circuits and have the same expressive power as AIGs, provided that AIGs allow for efficient detection of embedded XOR-subgraphs, which require specialized handling in some procedures. Therefore, our choice of AIGs is somewhat arbitrary and is motivated by its straight-forward interpretation and convenient implementation.

The material on windowing and reconvergence-driven cut computation in the following two subsections assumes a general Boolean network, but the synthesis algorithms in Section 3 are optimized assuming an AIG.

2.3 Windowing

Windowing is a method of limiting the scope of logic synthesis to work only on a small portion of a Boolean network. This method is indispensable for scalability when working with large Boolean networks arising in industrial applications.

The material in this section is adapted from [17], where windowing is used to compute don't-cares. A full presentation of the original algorithm is included because the windowing algorithms in this paper are based on it and extend it in a number of ways.

Definition. Two non-overlapping subsets of nodes, the *leaf set* and the *root set*, are in a leaf/root relation if every path from the PIs to any node in the root set passes through some node in the leaf set.

Definition. Given two subsets in the leaf/root relationship, its *window* is the subset of nodes of the network containing the root set together with all nodes on paths between the leaf set and the root set. The nodes in the leaf set are not included in the window.

Definition. A path connecting a pair of nodes is *distance-k* if it spans exactly k edges between the pair. Two nodes are *distance-k* from each other if the shortest path between them is distance-k.

The pseudo-code in Figure 2.3.1 and the example in Figure 2.3.2 describe the flow of a window construction algorithm. Procedure *Window* takes a node and two integers defining the number of logic levels on the fanin/fanout sides of the node to be included in the window. It returns the leaf set and the root set of the window.

```
(nodeset, nodeset) Window( node N, int nFanins, int nFanouts )
{
    nodeset I₁ = CollectNodesTFI( {N}, nFanins );
    nodeset O₁ = CollectNodesTFO( {N}, nFanouts );
    nodeset I₂ = CollectNodesTFI( O₁, nFanins + nFanouts );
    nodeset O₂ = CollectNodesTFO( I₁, nFanins + nFanouts);
    nodeset S = I₂ ∩ O₂;
    nodeset L = CollectLeaves( S );
    nodeset R = CollectRoots( S );
    return (L, R);
}
```

Figure 2.3.1. Computation of a window for a node.

The procedure *CollectNodesTFI* takes a set S of nodes and an integer $m \geq 0$, and returns a set of nodes on the fanin side that are distance-m or less from the nodes in S. An efficient implementation of this procedure for small m (for most applications, $m \leq 10$) iterates through the nodes that are distance-k ($0 \leq k \leq m$) from the given set. The distance-0 nodes are the original nodes. The distance-$(k+1)$ nodes are found by collecting the fanins of the distance-k nodes not visited before. The procedure *CollectNodesTFO* is similar.

Procedures *CollectLeaves* and *CollectRoots* take the set of the window's internal nodes and determine the leaves and roots of this window. The leaves are the nodes that do not belong to the given set but are fanins of at least one of the nodes in the set. The roots are the nodes that belong to the given set and are also fanins of at least one node not in the set. Note that some of the roots thus computed are not in the TFO cone of the original node, for which the window is being computed, and therefore can be dropped without violating the definition of the window and undermining the usefulness of the window for logic synthesis operations dealing with the node.
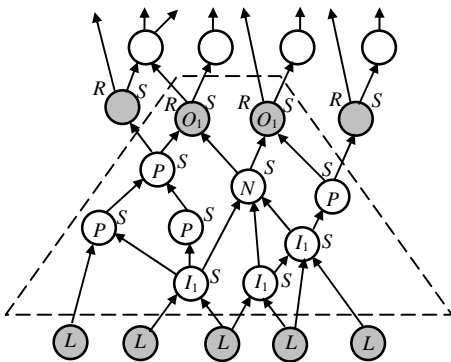

Figure 2.3.2. Example of the 1×1 window of node N.

We refer to the window constructed for a node by including n TFI logic levels and m TFO logic levels as an $n \times m$ window.

Example. Figure 2.3.2 shows a 1×1 window for node N in a network. The nodes labeled I_1, O_1, S, L, and R are in correspondence with the pseudo-code in Figure 2.3.1. The window's roots (top) and leaves (bottom) are shaded. Note that the nodes labeled by P do not belong to TFI and TFO of N, but represent reconvergent paths in the vicinity of N. The left-most root and right-most root are not in the TFO of N and can be dropped, as explained above.

2.4 Reconvergence-driven cut computation

Definition. *Reconvergence* occurs when the paths starting at the output of a node meet again before reaching the POs. Reconvergence is inevitable due to logic sharing in multi-level logic networks, but excessive reconvergence is often redundant.

Some applications, such as resubstitution or computation of a subset of satisfiability don't-cares of a node, require only one K-feasible cut of the node, of size between 5 and 20 inputs, depending on the computational effort allowed. Such a cut can be computed using procedure *CollectNodesTFI* presented in the previous section on windowing. For example, this procedure can be called for a node, resulting in a cut composed of the leaves of the TFI cone extending several logic levels down from the node.

However, a problem with this is that it is hard to predict how many logic levels to traverse to get a cut of the desired size. A second problem is that the cut computed this way may not lead to a good optimization because it includes few nodes or non-reconvergent (tree-like) logic structures. A small volume may lead to only a few nodes being available as resubstitution candidates. A tree-like structure does not lead to any don't-cares in the local space of the node. In both cases, the time spent computing the cut and attempting optimization using this cut would be wasted.

```
nodeset ReconvergenceDrivenCut( node N, int CutSizeLimit )
{
    nodeset Leaves = { N };
    nodeset Visited = { N };
    ConstructCut_rec( Leaves, Visited, CutSizeLimit );
    return Leaves;
}
ConstructCut_rec( nodeset Leaves, nodeset Visited, int CutSizeLimit )
{
    if ( Leaves contain only PI nodes )
        return;
    M = non-PI node in Leaves with the minimum LeafCost;
    if ( |Leaves| + LeafCost( M, Visited ) > CutSizeLimit )
        return;
    Leaves = Leaves ∪ fanins(M) \ M;
    Visited = Visited ∪ fanins(M);
    ConstructCut_rec( Leaves, Visited, CutSizeLimit );
}
int LeafCost( node M, nodeset Visited )
{
    int Cost = -1;
    for each fanin F of node M
        if F does not below to Visited
            Cost = Cost + 1;
    return Cost;
}
```

Figure 2.4.1. Reconvergence-driven cut computation.

In this section, we present a simple and efficient cut computation algorithm, which computes a cut close to a given size (if such cut exists) while heuristically maximizing the cut volume and the number of reconvergent paths subsumed in the cut. This algorithm

has been used recently in several applications, including technology-dependent resynthesis [19] and refactoring [20]. It will be used for resubstitution in Section 3.2.

Procedure *ReconvergenceDrivenCut* in Figure 2.4.1 computes one good-quality cut of a given size for node *N*. The procedure uses two sets of nodes to store the leaves and the visited nodes, and initializes both with {*N*}. Next, the recursive procedure *ConstructCut_rec* is called. On termination, the computed cut is in the leaf set.

Procedure *ConstructCut_rec* expands the cut by incrementally adding one node. If the cut is composed of only PIs, the procedure quits. Else, it selects a non-PI node *M* that minimizes the cost, equal to the number of nodes that will be added to the leaves if *M* is used to expand the cut. This cost is computed using procedure *LeafCost*. The cost can be -1 if the node and both of its fanins are currently in the cut. Otherwise, the cost is 0 or more. If the least expansion of the cut makes it exceed the cut-size limit, the procedure quits. Otherwise, it updates the leaves and the visited nodes and calls itself recursively.

The above procedure works by greedily minimizing the number of the cut leaves in each iteration. In doing so, it tends to subsume into the cut those nodes that contribute to the volume but not to the cut size. It also prefers nodes with the costs -1 and 0, which are the cut leaves with reconvergent paths inside the cut.

3 Logic synthesis algorithms for AIGs

In this section, we describe several algorithms for efficient AIG-based logic synthesis. The first algorithm (*rewriting*) was developed specifically for AIGs. Other algorithms (*resubstitution* and *redundancy removal*) are adapted for AIGs from traditional synthesis. Using AIGs simplifies these algorithms, improves their speed and scalability, and makes them easier to implement.

3.1 Rewriting

Rewriting is a fast greedy algorithm for minimizing the AIG size by iteratively selecting AIG subgraphs rooted at a node and replacing them with smaller pre-computed subgraphs, while preserving the functionality of the root node. Below we outline the rewriting algorithm [20] developed by extending [2] in the following ways:
- Using 4-input cuts instead of two-level subgraphs.
- Restricting rewriting to preserve the number of logic levels.
- Developing several variations of AIG rewriting that (a) look at larger subgraphs and (b) attempt to reduce the AIG depth.
- Experimental tune-up for logic synthesis applications.

For the purposes of AIG rewriting, all 4-feasible cuts of the nodes are found using the fast cut enumeration procedure [24][21]. For each cut, the Boolean function is computed and its NPN-class is determined by hash-table lookup. Fast manipulation of 4-variable functions is achieved by representing them using truth tables stored as 16-bit bit-strings. Altogether there are 222 NPN equivalence classes of 4-variable functions [23], of which only about 100 appear more than once as functions of 4-feasible cuts in the available benchmarks, and only about 40 of these have been found experimentally to lead to improvements in rewriting. The unifying characteristic of the useful NPN-classes of functions is that they are decomposable using simple disjoint-support decomposition [2].

All non-redundant AIG subgraphs of the representative functions of the useful equivalence classes are pre-computed in advance as a shared DAG with approximately 2,000 nodes and hashed by the truth table. This DAG is compiled into the program as an integer array, which noticeably reduced the setup time of the rewriting package.

```
Rewriting( network AIG, hash table PrecomputedStructures, bool UseZeroCost )
{
    for each node N in the AIG in the topological order {
        for each 4-input cut C of node N computed using cut enumeration {
            F = Boolean function of N in terms of the leaves of C
            PossibleStructures = HashTableLookup( PrecomputedStructures, F );
            // find the best logic structure for rewriting
            BestS = NULL; BestGain = -1;
            for each structure S in PossibleStructures {
                NodesSaved = DereferenceNode( AIG, N );
                NodesAdded = ReferenceNode( AIG, S );
                Gain = NodesSaved – NodesAdded;
                Dereference( AIG, S );
                Reference( AIG, N );
                if ( Gain > 0 || (Gain = 0 && UseZeroCost) )
                    if ( BestS = NULL || BestGain < Gain )
                        BestS = S; BestGain = Gain;
            }
            // use the best logic structure to update the netlist
            if ( BestS != NULL ) {
                NodesSaved = DereferenceNode( AIG, N );
                NodesAdded = ReferenceNode( AIG, S );
                assert( BestGain = NodesSaved – NodesAdded );
            }
        }
    }
}
```
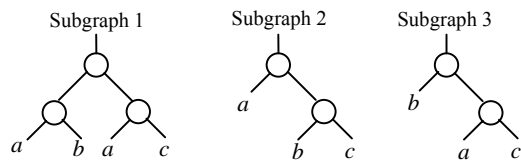
Figure 3.1.1. 4-input rewriting algorithm.

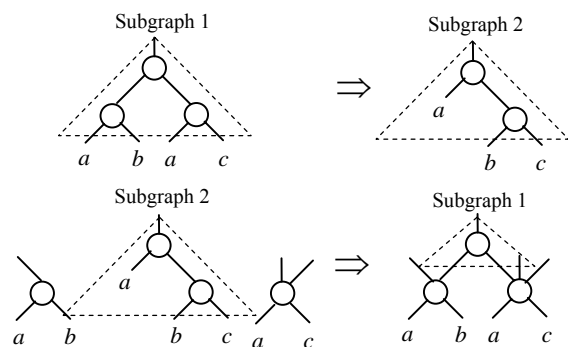**Figure 3.1.2**. Different AIG structures for function *F = abc*.

Figure 3.1.3. Two cases of AIG rewriting of a node.

Figure 3.1.1 shows the pseudo-code of the AIG rewriting procedure. The nodes are visited in the topological order. For each 4-input cut of a node, all pre-computed subgraphs are considered. Logic sharing between the new subgraphs and nodes already in the network is detected using an AIG with reference counters. For this, the old subgraph is de-referenced and the number of nodes, whose reference counts became 0, is returned. This is the number of nodes saved by not having the old subgraph in the network. Then, a new subgraph is added to the network while counting the number of new nodes and the nodes whose reference count changes from 0 to a

positive value. This is the increase in the number of nodes due to having the new subgraph in the network. The difference between the former and the latter numbers is the gain in the number of nodes if the replacement is done. The new node is de-referenced and the old node is referenced to return the AIG to its original state.

After trying all available subgraphs, the one that leads to the largest improvement at a node is used. If there is no improvement and "zero-cost replacement" is enabled, a new subgraph that does not increase the number of nodes is used.

Example. Figure 3.1.2 shows three AIGs for $F = abc$ that are pre-computed and stored. Figure 3.1.3 shows two instances of AIG rewriting. The upper part of the figure shows the situation when Subgraph 1 is detected and replaced by Subgraph 2. The lower part of the figure shows two nodes AND(a, b) and AND(a, c) that are already present in the network. In this case, Subgraph 2 can be replaced by Subgraph 1. In both cases, one node is reduced.

For further details on AIG rewriting and experimental comparison with the traditional logic synthesis, refer to [20].

3.2 Resubstitution

Resubstitution expresses the function of a node using other nodes (called *divisors*) already present in the network. The transformation is accepted if the new implementation of the node is in some sense better than its current implementation using the immediate fanins.

In the case of an AIG, the best outcome of resubstitution is when the whole MFFC of the node can be freed and the node's function can be expressed using a node that is currently outside of the node's MFFC. This outcome is called a 0-resubstitution because no new nodes are added to the AIG while the MFFC is removed. Similarly, if there is no 0-resubstitution, a 1-resubstitution exists if the function of the node can be expressed using two already present nodes and exactly one additional node.

This approach generalizes to *k-resubstitution*, which adds exactly k new nodes to the AIG and reduces the AIG size if the node's MFFC has at least $k+1$ nodes. This approach is conceptually similar to technology-dependent resynthesis based on resubstitution [15][19]. The difference is, in the case of AIGs, a technology-independent metric (the AIG size) is used, and for scalability and speed, the cuts with no more than 12-16 leaves are used. For such relatively small cuts, resubstitution can be performed without BDDs and SAT, using explicitly computed truth tables and exhaustive simulation.

```
Resubstitution( network AIG, int CutSizeLimit, int DivisorLimit, bool UseZeroCost )
{
    for each node n in the AIG in the topological order {
        int MffcSize = | MFFC( n )|;  assert( MffcSize ≥ 1 );
        nodeset C = ReconvergenceDrivenCut( n, CutSizeLimit );
        nodeset D = CollectNodesTFOChanged( C, Level(n), DivisorLimit );
        ComputeFunctions( D, n );
        nodeset R = Try0Resubstitution( n, D );
        if ( R == NULL && (MffcSize > 1 || (MffcSize == 1 && UseZeroCost)) )
            R = Try1Resubstitution( n, D );
        if ( R == NULL && (MffcSize > 2 || (MffcSize == 2 && UseZeroCost)) )
            R = Try2Resubstitution( n, D );
        if ( R == NULL && (MffcSize > 3 || (MffcSize == 3 && UseZeroCost)) )
            R = Try3Resubstitution( n, D );
        If ( R != NULL )
            UpdateNetwork( AIG, n, R );
    }
}
```

Figure 3.2.1. AIG resubstitution algorithm.

Figure 3.2.1 shows the pseudo-code of the resubstitution algorithm. For each node processed, the size of MFFC (containing at least the node itself) and a reconvergence-driven cut are computed as shown

in Section 2.4. Next, a limited TFO computation is found given the cut, the level of node n, and an upper bound on the number of divisors. This computation differs from *CollectNodesTFO* in Section 2.3 as follows: (a) instead of collecting the nodes that are no more than k levels from the starting set, the nodes are collected whose levels do not exceed the level of node n, (b) the leaves of the cut are collected but the nodes in the MFFC are skipped because they will be removed, and (c) there is a limit on the total number of nodes collected (*DivisorLimit*). Next, the Boolean functions of the collected nodes in the terms of the cut leaves are computed using exhaustive simulation.

Several resubstitutions of the node are attempted, each trying to add more nodes to the AIG. First, procedure *Try0Resubstitution* checks whether the Boolean function of node n is a constant or equal (up to complementation) to that of a divisor. If there is no such 0-resubstitution, *Try1Resubstitution* is tries to find two divisors that, when ANDed in some polarities, are equal (up to complementation) to the function of node n. Note that 1-resubsitution can only lead to improvement in the AIG size if the size of MFFC is more than one. Similar observations hold for the higher-order resubstitutions.

Our resource-aware implementation limits the set of divisors, D, because the complexity of k-resubstitution is $O(|D|^{k+1})$. We use $k = \{0, 1, 2, 3\}$. For this reason, the limit on the number of divisors computed in *CollectNodesTFOChanged* in Figure 3.2.1 is 150. Despite the high polynomial complexity of the resubstitution test, the total runtime is often dominated by collecting the devisors.
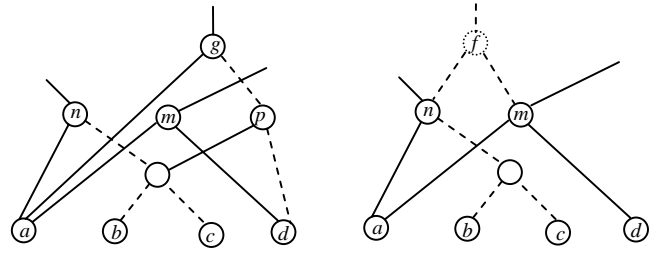
Figure 3.2.2. Example of AIG resubstitution.

Example. In the AIG shown in Figure 3.2.2 (left), node g has MFFC of size 2 composed of nodes g and p. One node can be saved if node g is replaced by the complement of node f, which does not belong to the original AIG but can be added on top of the already present nodes n and m, as shown in Figure 3.2.2 (right). Indeed, $g = a(b + c + d)$, $\overline{f} = n + m = a(b+c) + ad = a(b+c+d)$. This resubstitution is an example of the algebraic distributive transform.

3.3 Redundancy removal

Redundancy removal (RR) greedily removes network edges that do not contribute to the functionality of the network POs. RR is closely related to don't-care-based optimization of general Boolean networks, which simplifies the local functions of the nodes using flexibilities caused by the surrounding network. More than a decade of research resulted in several scalable computations for computing at a node, a subset of compatible observability don't-cares (CODCs) [25] or a subset of complete don't-cares using windowing [17].

These computations could be adapted to work on AIGs, but a don't-care at an input of a two-input AND node only can have no effect, or cause it to be replaced by a constant, an inverter, or a buffer. Thus any improvement achieved using don't-cares can be seen as due to a redundancy in the network when one or more fanin edges are redundant. Therefore, don't-care-based two-level minimization performed in [25][17] using ESPRESSO is not needed

for AIGs. This motivates the following efficient RR procedure based on structural analysis, simulation, and Boolean satisfiability.

First, the network is simulated to detect as many non-redundant edges as possible. The remaining edges are potentially redundant and are checked as shown in Figure 3.3.1. The procedure *RedundancyRemoval* iterates over candidate edges, creating a window for each edge. The window computation differs from the more general one presented in Section 2. During the computation of the TFO of the window leaves using procedure *CollectNodesTFO*, the target edge E is not traversed. As a result, the resulting TFO contains only the nodes reachable from the leaves through the paths bypassing E. If there is a path going from the fanout of E to a PO that does not pass through a node in the TFO, the redundancy of E cannot be proved using a window of the given size. Otherwise, it is possible that the information flow through E to the POs is blocked by reconvergent paths and E can be proved redundant. Procedure *CollectSubsetReachableFrom* finds the roots of the TFO reachable from the fanout of E. The TFI cone of these roots (*TFI2*) is collected. The intersection of *TFI2* and *TFO* gives the desired window. The combinational miter [5] is constructed, as shown in Figure 3.3.2.

```
RedundancyRemoval( network AIG, int WindowSize, int Timeout )
{
    edgeset E = candidate AIG edges not disproved by random simulation;
    for each candidate edge in E {
        W = CreateWindowForRR( E, WindowSize, WindowSize );
        if ( containing windows W of this size does not exist )
            continue;
        network Miter = ConstructMiter( W, E );
        if ( RandomSimulation( Miter ) ) // edge is not disproved by simulation
            if ( CheckRRusingSat( Miter, Timeout ) ) // proved redundant by SAT
                UpdateNetwork( AIG, E );
    }
}

(nodeset, nodeset) CreateWindowForRR( edge E, int nFanins, int nFanouts )
{
    nodeset TFI = CollectNodesTFI( {Fanin(E)}, nFanins );
    nodeset TFO = CollectNodesTFO( CollectLeaves(TFI), nFanins + nFanouts, E );
    if ( there is a path from Fanout(E) to a PO not passing through a node in TFO )
        return window of this size does not exist;
    nodeset Roots = CollectSubsetReachableFrom( CollectRoots(TFO), Fanout(E) );
    nodeset TFI2 = CollectNodesTFI( Roots, nFanins + nFanouts );
    nodeset S = TFI2 ∩ TFO;
    nodeset L = CollectLeaves( S );
    nodeset R = CollectRoots( S );
    return (L, R);
}
```

Figure 3.3.1. Redundancy removal for AIGs.

Random simulation is now applied to the miter for early detection of non-redundancy. Although this was performed at the beginning to filter out obvious non-redundant edges, additional simulation detects candidate redundant edges that became non-redundant through the intervening removal of some redundancies. If the edge remains a candidate, the miter is solved by SAT [22]. If the miter is unsat, the edge is redundant because the outputs of the original and modified windows are equal for all assignments to the leaves of the window. In this case, the redundant edge is removed from the AIG.

Example. In Figure 3.3.3, the AIG edge $g \rightarrow f$ is redundant because the function of h, the only fanout of f, does not change when the $g \rightarrow f$ is dropped (when g is replaced by constant 0). Indeed, $\overline{f} = ab + \overline{b}cde$, $h = \overline{f}\,bc = (ab + \overline{b}cde\,)bc = abc$. After the change, $h = \overline{f}\,bc = (ab+0)bc = abc$. Edge $g \rightarrow f$ can be proved redundant by considering a window with roots $\{h\}$ and leaves $\{a, b, c, d, e\}$.
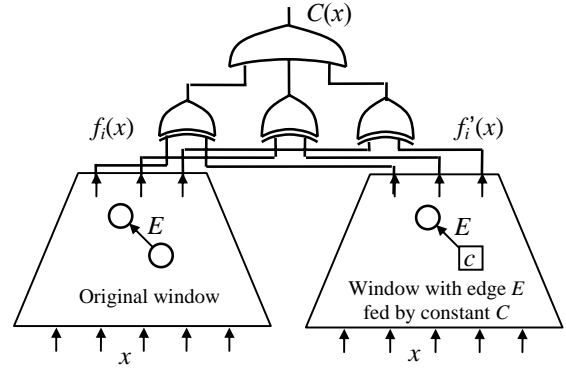

**Figure 3.3.2**. A miter used for RR.

Figure 3.3.3. Example of redundancy removal in AIGs.

An interesting aspect of RR in AIGs is that a fanin edge can be redundant with respect to one of the fanout edges. For example, consider a copy of a node for each fanout. A fanin edge may be redundant in one of the copies. Removing the redundancy in this copy is equivalent to replacing the associated fanout in the original circuit. Although the node count remains unchanged, RR with respect to a fanout removes false logic sharing leading to shorter delays and improved stuck-at testability.

4 Experimental Results

The presented algorithms were implemented in a public-domain logic synthesis and verification system ABC [1]. This section compares AIG rewriting with and without resubstitution for area-only AIG optimization, and reductions achieved by redundancy removal[1]. The subset of IWLS 2005 benchmarks [13] is selected for these experiments based on circuit size (to demonstrate scalability), and not on the improvements by the proposed algorithms. All results are verified using the equivalence checker in ABC [22].

Table 1 lists the benchmark names and the numbers of AIG nodes after an initial structural hashing (column *strash*), and then after using several scripts described in Figure 4.1:

- three iterations of rewriting (column **3×rw**),
- three iterations of resubstitution (column **3×rs**),
- three iterations of rewriting interleaved with resubstitution (column **3×rws**),
- an area-minimization script *compress*2 (column **com2**), and
- a script derived from *compress*2 by interleaving rewriting and resubstitution (column **com2rs**).

---

[1] For a comparison of AIG rewriting and traditional logic synthesis see [20].

3×rw: st; rw -l; rwz -l; rwz -l

3×rs: st; rs -K 6 -N 2 -l; rs -K 9 -N 2 -l; rs -K 12 -N 2 -l

3×rws: st; rw -l; rs -K 6 -N 2 -l; rwz -l; rs -K 9 -N 2 -l; rwz -l; rs -K 12 -N 2 -l

compress2: b -l; rw -l; rf -l; b -l; rw -l; rwz -l; b -l; rfz -l; rwz -l; b –l

compress2rs: b -l; rs -K 6 -l; rw -l; rs -K 6 -N 2 -l; rf -l; rs -K 8 -l; b -l; rs -K 8 -N 2 -l; rw -l; rs -K 10 -l; rwz -l; rs -K 10 -N 2 -l; b -l; rs -K 12 -l; rfz -l; rs -K 12 -N 2 -l; rwz -l; b -l

Semicolons separate individual commands in the scripts:

st is structural hashing

b is algebraic tree-balancing

rw is rewriting

rwz is rewriting with zero-cost replacement

rf is refactoring

rfz is refactoring with zero-cost replacement

rs is resubstitution

-K <num> is the limit on the number of cut inputs (default is 6)

-N <num> is the limit on the number of new nodes that can be added while performing resubstitution at each node (default is 1)

-l turns on the area-minimization mode (minimizing the number of AIG nodes while disregarding the number of AIG levels).

Figure 4.1. Description of various rewriting scripts.

Table 1 shows that, for area minimization, resubstitution alone cannot compete with rewriting (columns 3×**rw** and 3×**rs**). However, interleaving resubstitution with rewriting, leads to a 2-3% average improvement in the number of AIG nodes, compared to iterating rewriting (columns 3×**rw** and 3×**rws**). Also, using resubstitution as part of an area-minization script (**com2** versus **com2rs**) leads to an additional 2% improvement.

Table 1. The number of AIG nodes with and w/o resubstitution.

design	strash	3×rw	3×rs	3×rws	com2	com2rs
aes_core	21213	19814	20432	19413	19735	19326
des_perf	78299	69870	70577	66931	69196	65671
ethernet	19729	14415	18175	14153	13020	12784
pci_bridge32	22784	18062	22202	17899	17817	17701
usb_funct	15873	13779	14579	13299	13059	12540
vga_lcd	126711	91152	118926	90977	90844	90713
wb_conmax	47853	43847	45248	41666	40407	39645
average ratio	**1.21**	**1.00**	**1.13**	**0.97**	**0.96**	**0.94**

Table 2. The runtime (in seconds) with and w/o resubstitution.

design	strash	3×rw	3×rs	3×rws	com2	com2rs
aes_core	0.34	1.82	13.64	15.23	4.96	25.70
des_perf	1.00	16.31	34.19	47.16	31.89	73.42
ethernet	0.18	1.13	4.25	4.23	1.97	5.87
pci_bridge32	0.22	1.36	4.89	5.26	2.59	7.90
usb_funct	0.14	0.96	2.39	2.66	1.87	4.91
vga_lcd	1.71	20.50	120.75	92.98	31.00	131.53
wb_conmax	0.37	5.30	12.91	15.78	8.47	25.42
average ratio	**0.12**	**1.00**	**4.00**	**4.16**	**1.92**	**6.56**

It should be noted that although a 3% average reduction in area may not seem substantial, in practice, further reduction in the AIG size after the 10 passes of AIG rewriting, is hard to achieve. For example, even when a relatively expensive don't-care-based optimization is applied [17] and the resulting netlist is converted back into an AIG, the reduction is rarely more than 3%.

Table 2 shows runtimes measured on a 1.6GHz IBM ThinkPad with 1Gb RAM. It is clear that the current implementation of resubstitution is as scalable as AIG rewriting although several times slower. In most cases, the runtime of resubstitution was dominated by the computation of candidate Boolean divisors at each node, which uses backward reachability from the cut leaves (procedure

CollectNodesTFOChanged in Figure 3.2.1). On the other hand, the truth table computation using exhaustive simulation is relatively fast even for cuts with more than 10 inputs. Similarly, divisor checking, which is quadratic (cubic) in the number of divisors for resubstitution with 1 (2) additional nodes is less time-consuming because of the restriction on the number of divisors allowed.

Table 3. Results of applying redundancy removal.

design	strash	rr44	time, s	com2	rr44	time, s
aes_core	21213	20821	39.93	19735	19708	17.69
des_perf	78299	77300	84.08	69196	69097	67.86
ethernet	19729	18729	24.38	13020	12997	4.68
pci_bridge32	22784	21620	20.63	17817	17806	14.78
usb_funct	15873	15596	9.21	13059	13025	5.63
vga_lcd	126711	121508	295.47	90844	90837	75.68
wb_conmax	47853	47409	29.60	40407	40228	25.62
average ratio	**1.00**	**0.971**		**1.00**	**0.998**	

Table 3 shows early results of applying redundancy removal for all edges in the network using 4 x 4 windows. In the first section, RR (column **rr44**) was applied to the networks after structural hashing (column **strash**). In the second section, Script *compress2* (column **com2**) was applied before RR (column **rr44**). The runtimes of RR shown in columns **time** of Table 3 are preliminary because the current implementation is not optimized for runtime. In particular, we have not yet implemented early detection of non-redundant edges using random simulation and did not fine-tune the window computation for nodes with many fanouts.

It is surprising that RR cannot substantially reduce the networks after script *compress2*. This may speak for the efficiency of 10 passes of AIG rewriting. To verify that RR detects all redundancies we performed the following experiment for smaller benchmarks. We iterated RR until saturation followed by simplification using complete don't-cares [17] (with resubstitution disabled). Typically don't-care-based optimization could not further reduce the AIG size after RR had saturated, meaning that the current window-based implementation of RR indeed removes most of the redundancies.

5 Conclusions and Future Work

This paper presents several local transformations for combinational logic synthesis based on pre-computation of AIG subgraphs, reconvergence analysis, efficient bottom-up cut enumeration, Boolean matching, exhaustive simulation of logic cones of up to 16 inputs, and local resource-aware runs of Boolean satisfiability.

A distinctive feature of this work is that it uses AIGs exclusively during all synthesis steps as a simple multi-level logic representation. This allows local transformations to be applied at high speed and in a scalable manner. The cumulative gain of several rounds of local transformations compares in quality to the logic synthesis done in MVSIS and SIS, but is one or two orders of magnitude faster as well as applicable to larger examples [20].

The new AIG-based techniques have the potential for replacing traditional logic synthesis in CAD tools. Due to their extreme speed and good quality, they can be used in hardware emulation, early estimation of the design complexity, software synthesis [27], and preprocessing of miters before equivalence checking [22].

Future work will include extending the local transformations to use larger cut sizes. The challenge is to search an even larger space of alternative circuit structures while keeping runtime low, thereby allowing multiple optimization passes.

Acknowledgement

References

[1] Berkeley Logic Synthesis and Verification Group, *ABC: A system for sequential synthesis and verification, Release 60306*. http://www.eecs.berkeley.edu/~alanmi/abc/

[2] V. Bertacco and M. Damiani, "Disjunctive decomposition of logic functions," *Proc. ICCAD '97*, pp. 78-82.

[3] P. Bjesse and A. Boralv, "DAG-aware circuit compression for formal verification", *Proc. ICCAD '04*, pp. 42-49.

[4] K. S. Brace, R. L. Rudell, and R. E. Bryant, "Efficient implementation of a BDD package", *Proc. DAC '90*, pp. 40-45.

[5] D. Brand. "Verification of large synthesized designs". *Proc. ICCAD '93*, pp. 534 -537.

[6] R. Brayton and C. McMullen, "The decomposition and factorization of Boolean expressions," *Proc. ISCAS '82*, pp. 29-54.

[7] R. Brayton, G. Hachtel, A. Sangiovanni-Vincentelli, "Multilevel logic synthesis", *Proc. IEEE*, Vol. 78, Feb.1990.

[8] S. Chatterjee, A. Mishchenko, R. Brayton, X. Wang, and T. Kam, "Reducing structural bias in technology mapping", *Proc. ICCAD '05*. http://www.eecs.berkeley.edu/~alanmi/publications/2005/iccad05_map.pdf

[9] S. Chatterjee, A. Mishchenko, and R. Brayton, "Factor cuts", *Proc. IWLS '06*. ttp://www.eecs.berkeley.edu/~alanmi/publications/2006/iwls06_cut.pdf

[10] J. Cong, C. Wu and Y. Ding, "Cut ranking and pruning: Enabling a general and efficient FPGA mapping solution," *Proc. FPGA `99*, pp. 29-35.

[11] A. Darringer, W. H. Joyner, Jr., C. L. Berman, L. Trevillyan, "Logic synthesis through local transformations," *IBM J. of Research and Development*, Vol. 25(4), 1981, pp 272-280.

[12] L. Hellerman, "A catalog of three-variable Or-Inverter and And-Inverter logical circuits", *IEEE Trans. Electron. Comput.,* Vol. EC-12, June 1963, pp. 198-223.

[13] *IWLS '05 Benchmarks*. http://iwls.org/iwls2005/benchmarks.html

[14] K. Keutzer, "DAGON: Technology binding and local optimizations by DAG matching", *Proc. DAC '87*, pp. 617-623.

[15] V. N. Kravets and P. Kudva, "Implicit enumeration of structural changes in circuit optimization", *Proc. DAC '04*, pp. 438-441.

[16] A. Kuehlmann, "Dynamic transition relation simplification for bounded property checking". *Proc. ICCAD '04*, pp 50-57.

[17] A. Mishchenko and R. Brayton, "SAT-based complete don't-care computation for network optimization", *Proc. DATE '05*, pp. 418-423. http://www.eecs.berkeley.edu/~alanmi/publications/2005/date05_satdc.pdf

[18] A. Mishchenko, S. Chatterjee, R. Jiang, R. Brayton. "FRAIGs: A unifying representation for logic synthesis and verification". *ERL Technical Report, UC Berkeley, 2004*. http://www.eecs.berkeley.edu/~alanmi/publications/2005/tech05_fraigs.pdf

[19] A. Mishchenko, J. S. Zhang, S. Sinha, J. R. Burch, R. Brayton, and M. Chrzanowska-Jeske, "Using simulation and satisfiability to compute flexibilities in Boolean networks", *IEEE TCAD*, May 2006. http://www.eecs.berkeley.edu/~alanmi/publications/2005/tcad05_s&s.pdf

[20] A. Mishchenko, S. Chatterjee, and R. Brayton, "DAG-aware AIG rewriting: A fresh look at combinational logic synthesis", *Proc. DAC '06*. http://www.eecs.berkeley.edu/~alanmi/publications/2006/dac06_rwr.pdf

[21] A. Mishchenko, S. Chatterjee, and R. Brayton, "Improvements to technology mapping for LUT-based FPGAs", *Proc. FPGA'06*, pp. 41-49. http://www.eecs.berkeley.edu/~alanmi/publications/2006/fpga06_map.pdf

[22] A. Mishchenko, S. Chatterjee, R. Brayton, and N. Eén, "Improvements to combinational equivalence checking", *Proc. IWLS '06*. http://www.eecs.berkeley.edu/~alanmi/publications/2006/iwls06_cec.pdf

[23] S. Muroga, *Logic design and switching theory*, John Wiley & Sons, Inc., New York, NY, 1979

[24] P. Pan and C.-C. Lin, "A new retiming-based technology mapping algorithm for LUT-based FPGAs," *Proc. FPGA '98*, pp. 35-42.

[25] H. Savoj. *Don't cares in multi-level network optimization*. Ph.D. Dissertation, UC Berkeley, May 1992.

[26] E. Sentovich et al. "SIS: A system for sequential circuit synthesis." *Technical Report, UCB/ERI, M92/41, ERL*, Dept. of EECS, UC Berkeley, 1992.

[27] A. Solar-Lezama, R. Rabbah, R. Bodik, K. Ebcioglu, "Programming by sketching for bitstreaming programs", *ACM SIGPLAN Conference on Programming Language Design and Implementation* (*PLDI '05*), Chicago, IL, June 2005.